


A Mousetrap: Fooling Large Reasoning Models for Jailbreak with Chain of Iterative Chaos

Yang Yao, Xuan Tong, Ruofan Wang, Yixu Wang, Lujundong Li, Liang Liu, Yan Teng, Yingchun Wang

ABSTRACT

Large Reasoning Models (LRMs) have significantly advanced beyond traditional Large Language Models (LLMs) with their exceptional logical reasoning capabilities, yet these improvements introduce heightened safety risks. When subjected to jailbreak attacks, their ability to generate more targeted and organized content can lead to greater harm. Although some studies claim that reasoning enables safer LRM attacks, they overlook the inherent flaws within the reasoning process itself. To address this gap, we propose the first jailbreak attack targeting LRM, exploiting their unique vulnerabilities stemming from the advanced reasoning capabilities. Specifically, we introduce a Chaos Machine, a novel component to transform attack prompts with diverse one-to-one mappings. The chaos mappings iteratively generated by the machine are embedded into the reasoning chain, which strengthens the variability and complexity and also promotes a more robust attack. Based on this, we construct the Mousetrap framework, which makes attacks projected into nonlinear-like low sample spaces with mismatched generalization enhanced. Also, due to the more competing objectives, LRM gradually maintain the inertia of unpredictable iterative reasoning and fall into our trap. Success rates of the Mousetrap attacking o1-mini, claude-sonnet and gemini-thinking are as high as 96%, 86% and 98% respectively on our toxic dataset Trotter. On benchmarks such as AdvBench, StrongREJECT, and HarmBench, attacking claude-sonnet, well-known for its safety, Mousetrap can astonishingly achieve success rates of 87.5%, 86.58% and 93.13% respectively.

CONCEPT




ALGORITHM

Algorithm 1 Mousetrap attack

```

Input: dataset of PTQs;
Output: ASR and MLSs;
1: Make logs to record the result of PTQ
2: for PTQ in dataset:
3:   for length in [1,2,3]:
4:     succ_flag = 0
5:     Make logs to record the result of 3 attacks
6:     for equi_attack in range(3):
7:       DCPs, CTQ = ChaosMachine(PTQ, length)
8:       prompt = DCPs + CTQ
9:       response = AttackTarget(prompt)
10:      score = Judge(prompt, response)
11:      Record the PTQ result based on the score
12:      if all 3 times succeeded:
13:        succ_flag = 1
14:        Record the success with MLS
15:        break
16:      if succ_flag == 0:
17:        Record the failure
18: Calculate ASR
19: return ASR, MLSs
  
```


MOTIVATION



CONTRIBUTIONS

- First jailbreak attack targeting Large Reasoning Models (LRMs):** We present the first jailbreak attack specifically designed for LRM, addressing the overlooked vulnerabilities introduced by their advanced reasoning capabilities.
- Discovery of reasoning chain depth effect:** We observe that extending the reasoning chain in attack prompts enhances the success rate, revealing how LRM gradually maintain the inertia of unpredictable iterative reasoning and fall into our trap.
- Construction of a high-potency toxic benchmark:** We progressively identify the most toxic queries and construct Trotter, a highly potent benchmark containing three difficulty levels. This benchmark enables a more fine-grained evaluation of jailbreak performance and reveals the limitations of current safety alignment in LRM.
- Extensive experiments and strong results:** Our Mousetrap integrates the Chaos Machine with iterative reasoning chains to skillfully target the advanced reasoning abilities of LRM for jailbreaks. Notably, experiments conducted on nearly all well-known benchmarks in the field show that Mousetrap with a chain length of 3 consistently achieves success rates of no less than 86.58% even when attacking the famously safe Claude-3-5-Sonnet, which more compellingly reveals its robustness and generalizability.

EXPERIMENTS



RESOURCES RELEASED AT


GitHub.com/EvigByen/mousetrap

FIRST PREPRINT AT

arXiv.org/abs/2502.15806

AUTHOR'S HOMEPAGE

YAO' YANG
S PAGE
SINCE 2023




PRESENTED BY



Shanghai Artificial Intelligence Laboratory & The University of Hong Kong
& Fudan University & Hong Kong University of Science and Technology (Guangzhou)

FRAMEWORK

- Introduction of Chaos Machine and Mousetrap framework:** We propose Chaos Machine, a novel component that transforms attack prompts into diverse one-to-one mappings embedded within reasoning chains. Based on this, we construct Mousetrap, a jailbreak framework that projects attacks into nonlinear-like low sample spaces with mismatched generalization.



REGULARITY

For iterative reasoning attacks, as the chain length increases, the attack ability (opposite to safety alignment ability) rises, whereas the validity and correctness of reasoning decrease.

In practice, the expectation of effective attacks initially increases and then decreases. The horizontal position of the saddle point reflects the model's reasoning ability, while the vertical position corresponds to its safety alignment capability.

